

Multi-Broker Apache Kafka cluster in container:

- Run the zookeeper

```
./bin/zookeeper-server-start.sh config/zookeeper.properties
```

- Open a new tab and run the docker

```
docker run --rm \
-e BROKER_ID=0 \
-e ZOOKEEPER_CONNECT=host.docker.internal:2181 \
-e KAFKA_LISTENERS=PLAINTEXT://:9092 \
-e KAFKA_ADVERTISED_LISTENERS=PLAINTEXT://host.docker.internal:9092 \
-p 9092:9092 \
debezium/kafka:2.5
```

Part	Description
<code>docker run</code>	Starts a new Docker container.
<code>--rm</code>	Automatically removes the container after it stops. Useful for temporary containers.
<code>-e BROKER_ID=0</code>	Sets the Kafka broker ID to 0. Each broker in a cluster needs a unique ID.
<code>-e ZOOKEEPER_CONNECT=host.docker.internal:2181</code>	Tells Kafka where to find Zookeeper. Uses host machine's Zookeeper running on port 2181.

`-e KAFKA_LISTENERS=PLAINTEXT://:9092`

Kafka listens on port 9092 on all network interfaces using PLAINTEXT protocol.

`-e KAFKA_ADVERTISED_LISTENERS=PLAINTEXT://host.docker.internal:9092`

Kafka tells clients to connect using this advertised address (important for accessing Kafka from the host).

`-p 9092:9092`

Maps container's Kafka port 9092 to host's port 9092, allowing external access.

`debezium/kafka:2.5`

Specifies the Kafka Docker image from Debezium, version 2.5. (If you use `:latest` and it doesn't exist, it throws an error.)

But with newer versions of KAFKA, it prefers mentioning `NODE_ID` over broker ID
Create 3 brokers:

```
docker run --rm \
-e NODE_ID=0 \
-e ZOOKEEPER_CONNECT=host.docker.internal:2181 \
-e KAFKA_LISTENERS=PLAINTEXT://:9092 \
-e KAFKA_ADVERTISED_LISTENERS=PLAINTEXT://host.docker.internal:9092 \
-p 9092:9092 \
debezium/kafka:2.5
```

```
docker run --rm \
-e NODE_ID=1 \
-e ZOOKEEPER_CONNECT=host.docker.internal:2181 \
-e KAFKA_LISTENERS=PLAINTEXT://:9093 \
-e KAFKA_ADVERTISED_LISTENERS=PLAINTEXT://host.docker.internal:9093 \
-p 9093:9093 \
debezium/kafka:2.5
```

```
docker run --rm \
-e NODE_ID=2 \
-e ZOOKEEPER_CONNECT=host.docker.internal:2181 \
-e KAFKA_LISTENERS=PLAINTEXT://:9094 \
-e KAFKA_ADVERTISED_LISTENERS=PLAINTEXT://host.docker.internal:9094 \
-p 9094:9094 \
debezium/kafka:2.5
```

Check if the three brokers are running by using the command: `docker ps`

```
Last login: Fri Jun 13 13:29:16 on ttys003
(base) poojamanjunatha@poojas-air kafka_2.13-3.9.1 % docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS
600cdb9953bc	debezium/kafka:2.5	"/docker-entrypoint... cp, 0.0.0.0:9094->9094/tcp	9 minutes ago	Up 9 minutes	9092/t
be4fd9b7f29e	6551666480a4	"/docker-entrypoint... k8s_kafka_kafka-0_default_dd4dc3bb-9622-4fa2-82b4-bf89691205d1_8	10 minutes ago	Up 10 minutes	
5f871fc872d0	debezium/kafka:2.5	"/docker-entrypoint... 0:9092->9092/tcp	11 minutes ago	Up 11 minutes	0.0.0.

Next step is to run commands inside the container. The command is:

```
docker exec -it 600cdb9953bc /bin/bash
```

It gives you **interactive shell access** inside the container with ID `600cdb9953bc`. This is just like SSH'ing into a VM.

You'll be dropped into the container's terminal, where you can:

- Explore the filesystem
- Check logs
- Run commands (e.g., `kafka-topics.sh`)
- Troubleshoot issues

Create a topic called `mytopic` in the broker with the replication factor of 3 as there are 3 brokers.

```
./bin/kafka-topics.sh --create --topic mytopic --bootstrap-server host.docker.internal:9094 --partitions 1 --replication-factor 3
```

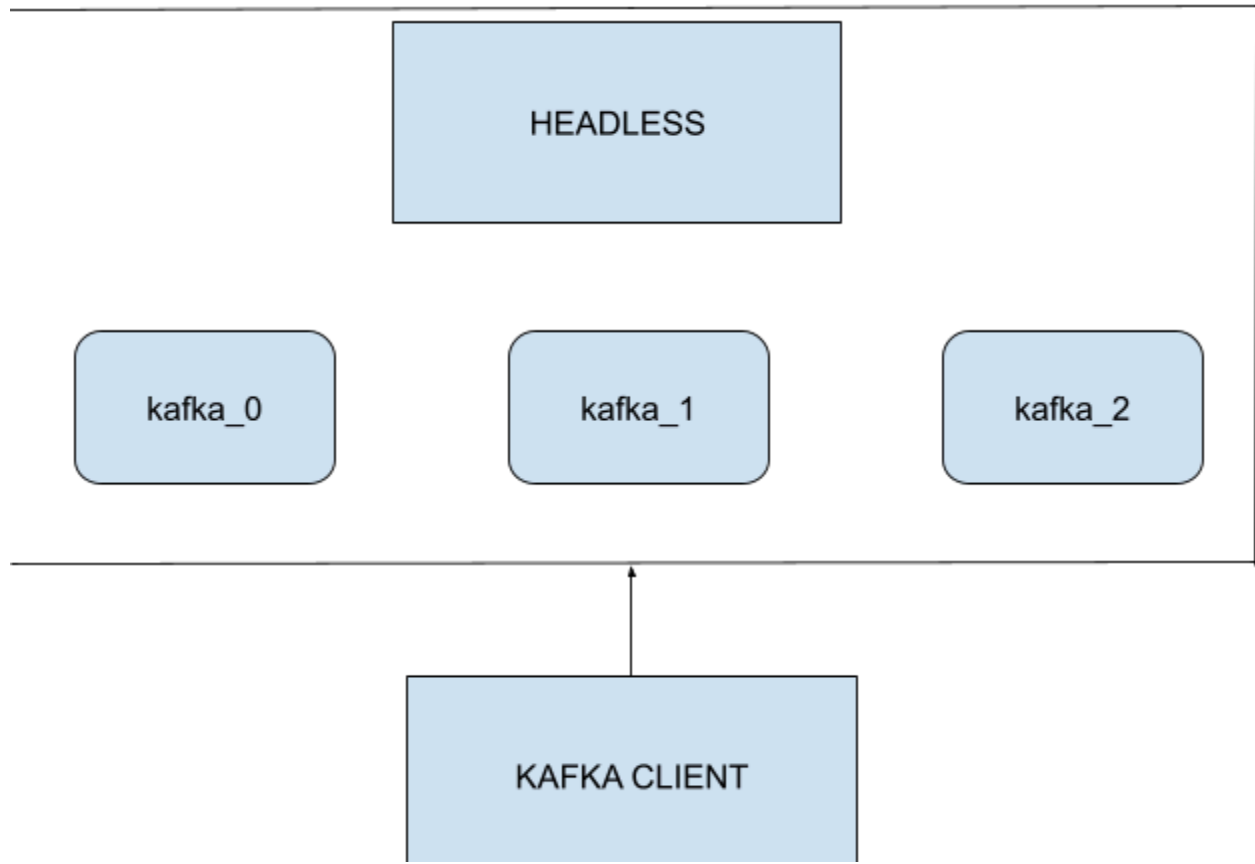
Produce messages in this topic using the command:

```
./bin/kafka-console-producer.sh --topic mytopic --bootstrap-server host.docker.internal:9094
```

```

--version
messages to.
Display Kafka version.
[kafka@600cdb9953bc ~]$ ./bin/kafka-console-producer.sh --topic mytopic --bootstrap-server host.docker.internal:9094
[>Hello
[>Bye
[>

```



These are **Kafka broker pods** managed as part of a **StatefulSet** in Kubernetes. They are named sequentially:

Pod Name	Broker ID
----------	-----------

kafka-0	0
---------	---

kafka-1	1
---------	---

kafka-2	2
---------	---

Each pod runs a Kafka broker and:

- Stores and serves partitions.
- Works with ZooKeeper (or Kraft mode in newer versions).
- Maintains state (thanks to StatefulSet: each pod gets a persistent identity and volume).

A **headless service** in Kubernetes is a service **without a cluster IP** (`clusterIP: None`). It is used for **direct pod-to-pod communication**.

Why use a headless service for Kafka?

- Brokers need to **know and directly communicate** with each other.
- Clients need to resolve **each broker individually** (e.g., `kafka-0.kafka-headless.default.svc.cluster.local`).
- This is **essential** for high availability and partition leadership.

So `kafka-headless` enables:

the **Kafka client** (producer or consumer) is usually:

- A separate **Kubernetes pod** or microservice.
- Connects to the Kafka cluster using the broker DNS names exposed by the **headless service**.
- Example bootstrap server config:

What are node ports?

Until now, we have accessed the broker from inside the container. But how do we access it from the outside? With the help of node ports: **NodePort** is a type of Kubernetes **Service** that exposes a pod **to the outside world (external clients)** by opening a specific port on **every node in the cluster**

Concept	Description
NodePort	A port (usually in the range <code>30000–32767</code>) exposed on every node in your cluster.
Access	You can access the service using <code>NodeIP:NodePort</code> .

Target It forwards traffic to the **ClusterIP service**, which then routes it to the appropriate pod(s).

To access the docker from outside:

1. Make sure zookeeper is running
2. Docker and Kubernetes must be running

Run the commands:

```
kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
docker-desktop	Ready	control-plane	12d	v1.32.2

Create a file called headless-service-kafka.yaml file and paste the content and then run the command:

```
kubectl apply -f headless-service-kafka.yaml
```

```
EOF
(base) poojamanjunatha@poojas-air kafka_2.13-3.9.1 % kubectl apply -f headless-service-kafka.yaml
service/kafka created
```

To check if it has been created:

Run the command

Kubectl get service

```
(base) poojamanjunatha@poojas-air kafka_2.13-3.9.1 % kubectl apply -f headless-service-kafka.yaml
service/kafka created
(base) poojamanjunatha@poojas-air kafka_2.13-3.9.1 % kubectl get service
NAME          TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)    AGE
kafka         ClusterIP   None         <none>         9092/TCP    2m48s
kubernetes    ClusterIP   10.96.0.1    <none>         443/TCP     12d
(base) poojamanjunatha@poojas-air kafka_2.13-3.9.1 %
```

Create a file for bootstrap-service-kafka as well and apply it .

Create 3 node ports file :

```
[EOF]
(base) poojamanjunatha@poojas-air kafka_2.13-3.9.1 % kubectl apply -f kafka-0-service.yaml -f kafka-1-service.yaml -f kafka-2-service.yaml

service/kafka-0 created
service/kafka-1 created
service/kafka-2 created
(base) poojamanjunatha@poojas-air kafka_2.13-3.9.1 %
```

```
base) poojamanjunatha@poojas-air kafka_2.13-3.9.1 % kubectl get service
NAME          TYPE          CLUSTER-IP    EXTERNAL-IP    PORT(S)          AGE
afka           ClusterIP     None           <none>          9092/TCP          30m
afka-0         NodePort      10.101.193.52 <none>          9093:30000/TCP    4m29s
afka-1         NodePort      10.108.109.51 <none>          9093:30001/TCP    4m29s
afka-2         NodePort      10.100.249.21 <none>          9093:30002/TCP    4m29s
afka-bootstrap NodePort      10.110.186.14 <none>          9092:30003/TCP    14m
kubernetes     ClusterIP     10.96.0.1     <none>          443/TCP           12d
base) poojamanjunatha@poojas-air kafka_2.13-3.9.1 %
```

Create and apply another file called statefulset-multi-broker.yaml and apply it

Run: `kubectl get pods -w` and this should be seen

Produce and consume messages